

SVM Homework


Megha Pandit

March 13, 2019

Problem 3

(a)

```
x1 <- c(3, 2, 4, 1, 2, 4, 4)
x2 <- c(4, 2, 4, 4, 1, 3, 1)
cols <- c("red", "red", "red", "red", "blue", "blue", "blue")
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
```



(b)


Since the two classes are very distinctly separable, we can see from the plot that the hyperplane must lie in between the points $\{(2,2), (4,4)\}$ and $\{(2,1), (4,3)\}$. Therefore, the hyperplane will pass through the points $(2, 1.5)$ and $(4, 3.5)$. The equation for the line passing through these

two points is:

$$x_2 = -0.5 + x_1$$

Therefore, the intercept and slope are -0.5 and 1 respectively.

```
#Optimal separating hyperplane  
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))  
abline(-0.5, 1)
```




(c)

The classification rule here would be: Classify as Red if $x_2 - x_1 + 0.5 > 0$, and, classify as Blue if $x_2 - x_1 + 0.5 < 0$

(d)

```
#Margin for maximal margin hyperplane  
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))  
abline(-0.5, 1)  
abline(-1, 1, lty = 2)  
abline(0, 1, lty = 2)
```



(e)

The support vectors for the maximal margin classifier are the points (2,1), (2,2), (4,3) and (4,4).


(f)

Since the 7th observation is not a support vector, a slight change in its position will not affect the maximal margin hyperplane. This is evident even from the above plot.

(g)


The equation $x_2 = -0.25 + x_1$ will also separate all the observations but is not an optimal hyperplane because the margin is smaller than the optimal option.

```
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
abline(-0.25, 1)
abline(0, 1, lty = 2)
abline(-0.5, 1, lty = 2)
```



(h)

```
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
abline(-0.5, 1)
abline(-1, 1, lty = 2)
abline(0, 1, lty = 2)
points(3,1, col = "red")
```



From the above plot, we see that after adding the new point, the hyperplane cannot separate the classes.


Problem 5

(a)

```
set.seed(9)
x1 <- runif(500) - 0.5
x2 <- runif(500) - 0.5
y <- 1*(x1^2 - x2^2 > 0)
```

(b)

```
plot(x1, x2, col = ifelse(y, "red", "blue"))
```




(c)

```
#Logistic regression
df <- data.frame(x1, x2, y)
fit_glm <- glm(y ~ x1 + x2, data = df, family = binomial)
fit_glm

##
## Call:  glm(formula = y ~ x1 + x2, family = binomial, data = df)
##
## Coefficients:
## (Intercept)          x1          x2
##    0.05514      0.38587     -0.02653
##
## Degrees of Freedom: 499 Total (i.e. Null);  497 Residual
## Null Deviance:      692.8
## Residual Deviance: 691.2    AIC: 697.2
```

(d)

```
pred_fit <- predict(fit_glm, data.frame(x1,x2))
plot(x1, x2, col = ifelse(pred_fit > 0, "red", "blue"), pch = ifelse(as.integer(pred_fit > 0) == y, 1, 4))
```



In the above plot, the circles are the observations that have been classified correctly and the crosses are the ones that are misclassified. The decision boundary looks linear.

(e)

```
fit_glm1 <- glm(y ~ poly(x1, 2) + poly(x2, 2), data = df, family = binomial)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit_glm1)
```

```
##
```

```
## Call:
```

```
## glm(formula = y ~ poly(x1, 2) + poly(x2, 2), family = binomial,  
##      data = df)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min      1Q      Median      3Q      Max  
## -1.079e-03 -2.000e-08  2.000e-08  2.000e-08  9.076e-04
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)      43.78    3063.63   0.014   0.989  
## poly(x1, 2)1    1360.39  102905.10   0.013   0.989
```

```
## poly(x1, 2)2 21374.91 785951.63 0.027 0.978
## poly(x2, 2)1 -119.10 88918.85 -0.001 0.999
## poly(x2, 2)2 -21333.50 788724.67 -0.027 0.978
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6.9276e+02 on 499 degrees of freedom
## Residual deviance: 2.4730e-06 on 495 degrees of freedom
## AIC: 10
##
## Number of Fisher Scoring iterations: 25


fit_glm2 <- glm(y ~ x1 + x2 + x1*x2, data = df, family = binomial)
summary(fit_glm2)

##
## Call:
## glm(formula = y ~ x1 + x2 + x1 * x2, family = binomial, data = df)
##
## Deviance Residuals:
## Min 1Q Median 3Q Max
## -1.342 -1.199 1.050 1.144 1.291
##
## Coefficients:
## Estimate Std. Error z value Pr(>|z|)
## (Intercept) 0.05206 0.08983 0.580 0.562
## x1 0.38234 0.31036 1.232 0.218
## x2 -0.01969 0.31760 -0.062 0.951
## x1:x2 0.64537 1.12041 0.576 0.565
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 692.76 on 499 degrees of freedom
## Residual deviance: 690.87 on 496 degrees of freedom
## AIC: 698.87
##
## Number of Fisher Scoring iterations: 3

None of the coefficient estimates are statistically significant.
```

(f)


```
pred_fit1 <- predict(fit_glm1, df)
plot(x1, x2, col = ifelse(pred_fit1 > 0, "red", "blue"), pch = ifelse(as.integer(pred_fit1 > 0) == y, 1,
```

The plot above shows a non-linear decision boundary and all the observations are correctly classified. Also, the decision boundary is similar to the true decision boundary.

(g)

```
#Support Vector Classifier
df$y <- as.factor(df$y)
fit_svc <- svm(y ~ x1 + x2, data = df, kernel = "linear")
pred_svc <- predict(fit_svc, df, type = "response")
plot(x1, x2, col = ifelse(pred_svc != 0, "red", "blue"), pch = ifelse(pred_svc == y, 1, 4))
```




In the above plot, the circles represent observations that have been classified correctly and crosses represent the observations that have been misclassified.

(h)

```
#SVM with non-linear kernel

fit_svm <- svm(y ~ x1 + x2, data = df, kernel = "polynomial", degree = 2)
pred_svm <- predict(fit_svm, df, type = "response")
plot(x1, x2, col = ifelse(pred_svm != 0, "red", "blue"), pch = ifelse(pred_svm == y, 1, 4))
```



There is a drastic improvement in classification compared to the linear kernel.

(i)

From the results, we see that SVM with a polynomial kernel of degree 2 performs quite well but still misclassifies some observations. In contrast, logistic regression with non-linear functions of predictors (polynomials of degree 2) does not misclassify any observations and definitely performs better than SVM.

Problem 7

(a)


```
data("Auto")
Auto$Y <- ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
Auto$Y <- as.factor(Auto$Y)
```

(b)

```
set.seed(9)
cost <- data.frame(cost = seq(0.01, 100, length.out = 10))
svm_tune <- tune(svm, Y ~ ., data = Auto, kernel = "linear", ranges = cost)
summary(svm_tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 11.12
##
## - best performance: 0.02557692
##
## - Detailed performance results:
##      cost      error dispersion
## 1    0.01 0.07679487 0.04850079
## 2   11.12 0.02557692 0.02417544
## 3   22.23 0.03070513 0.02649650
## 4   33.34 0.03326923 0.03211170
## 5   44.45 0.03326923 0.03211170
## 6   55.56 0.03326923 0.03211170
## 7   66.67 0.03326923 0.03211170
## 8   77.78 0.03326923 0.03211170
## 9   88.89 0.03326923 0.03211170
## 10 100.00 0.03326923 0.03211170

plot(svm_tune$performances[,c(1,2)], type = "l")
```



Cost = 11.12 seems to perform the best. This is also seen in the plot.

(c)

```
#Polynomial Kernel
para <- data.frame(cost = seq(0.01, 100, length.out = 5), degree = seq(1, 100, length.out = 5))

svm_poly <- tune(svm, Y ~ ., data = Auto, kernel = "polynomial", ranges = para)
summary(svm_poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##   100      1
##
## - best performance: 0.01519231
##
## - Detailed performance results:
##       cost degree      error dispersion
## 1    0.0100    1.00 0.56115385 0.02988888
```

```
## 2 25.0075 1.00 0.05089744 0.03945905
## 3 50.0050 1.00 0.02782051 0.04182274
## 4 75.0025 1.00 0.02025641 0.02601613
## 5 100.0000 1.00 0.01519231 0.02135480
## 6 0.0100 25.75 0.56115385 0.02988888
## 7 25.0075 25.75 0.56115385 0.02988888
## 8 50.0050 25.75 0.56115385 0.02988888
## 9 75.0025 25.75 0.56115385 0.02988888
## 10 100.0000 25.75 0.56115385 0.02988888
## 11 0.0100 50.50 0.56115385 0.02988888
## 12 25.0075 50.50 0.56115385 0.02988888
## 13 50.0050 50.50 0.56115385 0.02988888
## 14 75.0025 50.50 0.56115385 0.02988888
## 15 100.0000 50.50 0.56115385 0.02988888
## 16 0.0100 75.25 0.56115385 0.02988888
## 17 25.0075 75.25 0.56115385 0.02988888
## 18 50.0050 75.25 0.56115385 0.02988888
## 19 75.0025 75.25 0.56115385 0.02988888
## 20 100.0000 75.25 0.56115385 0.02988888
## 21 0.0100 100.00 0.56115385 0.02988888
## 22 25.0075 100.00 0.56115385 0.02988888
## 23 50.0050 100.00 0.56115385 0.02988888
## 24 75.0025 100.00 0.56115385 0.02988888
## 25 100.0000 100.00 0.56115385 0.02988888
```

Cost of 100 with degree 1 seems to perform the best.

#Radial Kernel

```
params <- data.frame(cost=seq(0.01,100,length.out = 5),gamma=seq(0.1,100,length.out = 5))
svm_radial <- tune(svm, Y ~ ., data = Auto, kernel = "radial", ranges = params)
summary(svm_radial)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
## 25.0075 0.1
##
## - best performance: 0.02532051
##
## - Detailed performance results:
##   cost gamma error dispersion
## 1 0.0100 0.100 0.19365385 0.08464044
## 2 25.0075 0.100 0.02532051 0.03129777
## 3 50.0050 0.100 0.02532051 0.03355075
## 4 75.0025 0.100 0.02532051 0.03355075
## 5 100.0000 0.100 0.02532051 0.03355075
## 6 0.0100 25.075 0.56115385 0.03834765
## 7 25.0075 25.075 0.54596154 0.03688365
## 8 50.0050 25.075 0.54596154 0.03688365
## 9 75.0025 25.075 0.54596154 0.03688365
```


```
## 10 100.0000 25.075 0.54596154 0.03688365
## 11 0.0100 50.050 0.56115385 0.03834765
## 12 25.0075 50.050 0.55608974 0.03870442
## 13 50.0050 50.050 0.55608974 0.03870442
## 14 75.0025 50.050 0.55608974 0.03870442
## 15 100.0000 50.050 0.55608974 0.03870442
## 16 0.0100 75.025 0.56115385 0.03834765
## 17 25.0075 75.025 0.56115385 0.03834765
## 18 50.0050 75.025 0.56115385 0.03834765
## 19 75.0025 75.025 0.56115385 0.03834765
## 20 100.0000 75.025 0.56115385 0.03834765
## 21 0.0100 100.000 0.56115385 0.03834765
## 22 25.0075 100.000 0.56115385 0.03834765
## 23 50.0050 100.000 0.56115385 0.03834765
## 24 75.0025 100.000 0.56115385 0.03834765
## 25 100.0000 100.000 0.56115385 0.03834765
```

(d)

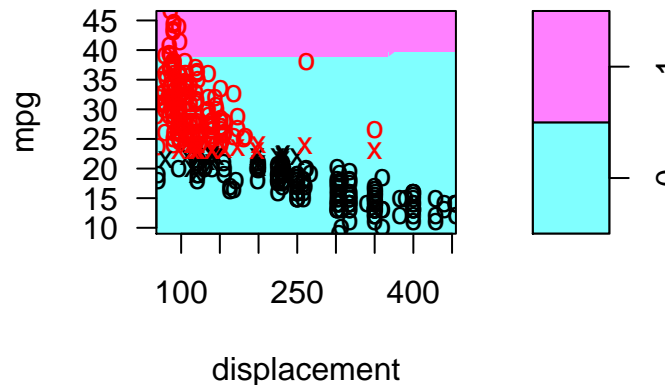
```
linear <- svm(Y ~ ., data = Auto, kernel = "linear", cost = 11.12)
polynomial <- svm(Y ~ ., data = Auto, kernel = "polynomial", cost = 100, degree = 1)
radial <- svm(Y ~ ., data = Auto, kernel = "radial", cost = 25.0075, gamma = 0.1)
pair_plot <- function(a){
  for (name in names(Auto)[!(names(Auto) %in% c("mpg", "Y", "name"))])
    plot(a, Auto, as.formula(paste("mpg~", name, sep = "")))
}

pair_plot(linear)
```


SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo



SVM classification plo



SVM classification plo




tion plots for linear kernel.


The above are the SVM classifica-

```
pair_plot(polynomial)
```


SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo



SVM classification plo




SVM classification plo




Thde above plots are the SVM classification plots for polynomial kernel.

```
pair_plot(radial)
```

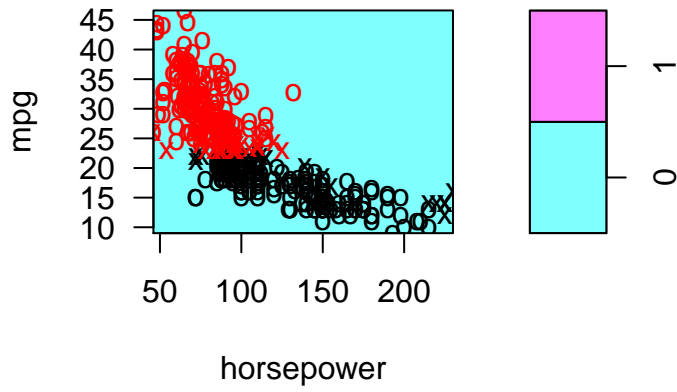
SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo



SVM classification plo



SVM classification plo



Problem 8

(a)

```
data("OJ")
set.seed(9)
train_oj <- sample(nrow(OJ), 800)
oj_train <- OJ[train_oj,]
oj_test <- OJ[-train_oj,]
```

(b)

```
oj_svc <- svm(Purchase ~ ., data = oj_train, kernel = "linear", cost = 0.01)
summary(oj_svc)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##   gamma:  0.05555556
##
## Number of Support Vectors:  427
```

```
##
## ( 214 213 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

The support vector classifier creates 432 support vectors out of the 800 training observations. Out of the 432 support vectors, 214 belong to level CH and 213 to level MM.

(c)

```
#Training error rate

pred_train <- predict(oj_svc, oj_train)
table(pred_train, oj_train$Purchase)
```

```
##
## pred_train CH MM
##          CH 428 70
##          MM  61 241
```

```
#Test error rate

pred_test <- predict(oj_svc, oj_test)
table(pred_test, oj_test$Purchase)
```

```
##
## pred_test CH MM
##          CH 143 33
##          MM  21 73
```

```
(tr_error<- (70+61)/(428+70+61+241))
```

```
## [1] 0.16375
```

```
(te_error <- (33+21)/(143+33+21+73))
```

```
## [1] 0.2
```

The training error rate is 16.37% and the test error rate is 20%.

(d)

```
#For optimal cost

oj_tune <- tune(svm, Purchase ~ ., data = oj_train, kernel = "linear", ranges =
               data.frame(cost = seq(0.01, 10, length.out = 25)))
summary(oj_tune)
```

```
##
## Parameter tuning of 'svm':
```

```
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 1.25875
##
## - best performance: 0.1575
##
## - Detailed performance results:
##   cost  error dispersion
## 1  0.01000 0.16750 0.05075814
## 2  0.42625 0.16250 0.04082483
## 3  0.84250 0.16250 0.04039733
## 4  1.25875 0.15750 0.03446012
## 5  1.67500 0.15750 0.03827895
## 6  2.09125 0.15750 0.04090979
## 7  2.50750 0.16125 0.04427267
## 8  2.92375 0.16125 0.04427267
## 9  3.34000 0.16250 0.04289846
## 10 3.75625 0.16375 0.04387878
## 11 4.17250 0.16375 0.04387878
## 12 4.58875 0.16250 0.04602234
## 13 5.00500 0.16500 0.04518481
## 14 5.42125 0.16500 0.04779877
## 15 5.83750 0.16500 0.04779877
## 16 6.25375 0.16625 0.04931827
## 17 6.67000 0.16625 0.04931827
## 18 7.08625 0.16875 0.05245699
## 19 7.50250 0.17000 0.05041494
## 20 7.91875 0.16875 0.04973890
## 21 8.33500 0.16875 0.04973890
## 22 8.75125 0.17000 0.04794383
## 23 9.16750 0.16750 0.05210833
## 24 9.58375 0.16750 0.05210833
## 25 10.00000 0.17000 0.04972145
```

From the results, we see that the optimal cost is **3.75625**

(e)

```
#Training error rate
oj_svm <- svm(Purchase ~ ., data = oj_train, kernel = "linear", cost = oj_tune$best.parameters$cost)
svm_train <- predict(oj_svm, oj_train)
table(svm_train, oj_train$Purchase)

##
## svm_train  CH  MM
##           CH 426  58
##           MM  63 253

(tr_err <- (58+64)/(425+58+64+253))
```

```
## [1] 0.1525
#Test error rate
svm_test <- predict(oj_svm, oj_test)
table(svm_test, oj_test$Purchase)
```

```
##
## svm_test  CH  MM
##          CH 142 27
##          MM  22 79
(te_err <- (27+22)/(142+27+22+79))
```

```
## [1] 0.1814815
```

The training error rate is 15.25% and the test error rate is 18.15%.

(f)

Radial Kernel

```
oj_radial <- svm(Purchase ~ ., data = oj_train, kernel = "radial")
summary(oj_radial)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:    1
##   gamma:    0.05555556
##
## Number of Support Vectors:  362
##
## ( 184 178 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

The SVM with radial kernel creates 624 support vectors out of the 800 training observations. Out of the 624 support vectors, 313 belong to level CH and 311 to level MM.

```
#Training error rate
radial_train <- predict(oj_radial, oj_train)
table(radial_train, oj_train$Purchase)
```

```
##
## radial_train  CH  MM
```



```
##          CH 441  70
##          MM  48 241
```

```
#Test error rate
```

```
radial_test <- predict(oj_radial, oj_test)
table(radial_test, oj_test$Purchase)
```

```
##
## radial_test  CH  MM
##           CH 147  34
##           MM  17  72
```

The training error rate is 14.75% and the test error rate is 18.89%.

```
#For optimal cost
```

```
radial_tune <- tune(svm, Purchase ~ ., data = oj_train, kernel = "radial", ranges =
                    data.frame(cost = seq(0.01, 10, length.out = 25)))
summary(radial_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 1.25875
##
## - best performance: 0.17375
##
## - Detailed performance results:
##   cost  error dispersion
## 1  0.01000 0.38875 0.04910660
## 2  0.42625 0.17500 0.03996526
## 3  0.84250 0.17625 0.03701070
## 4  1.25875 0.17375 0.04143687
## 5  1.67500 0.17875 0.04291869
## 6  2.09125 0.17875 0.04210189
## 7  2.50750 0.17875 0.03955042
## 8  2.92375 0.18375 0.04332131
## 9  3.34000 0.18250 0.04216370
## 10 3.75625 0.18250 0.04216370
## 11 4.17250 0.18125 0.04458528
## 12 4.58875 0.18250 0.04456581
## 13 5.00500 0.18375 0.04678927
## 14 5.42125 0.18500 0.04923018
## 15 5.83750 0.18625 0.04910660
## 16 6.25375 0.18750 0.04787136
## 17 6.67000 0.18875 0.04730589
## 18 7.08625 0.19125 0.04641674
## 19 7.50250 0.18875 0.04581439
## 20 7.91875 0.18875 0.04581439
## 21 8.33500 0.19000 0.04706674
## 22 8.75125 0.19250 0.04571956
```

```
## 23  9.16750 0.19250 0.04571956
## 24  9.58375 0.19125 0.04332131
## 25 10.00000 0.19000 0.04199868
```

The optimal cost is 0.42625.

#Training error rate

```
radial_svm <- svm(Purchase ~ ., data = oj_train, kernel = "radial", cost = radial_tune$best.parameters$
svm_rad <- predict(radial_svm, oj_train)
table(svm_rad, oj_train$Purchase)
```

```
##
## svm_rad  CH  MM
##          CH 444  70
##          MM  45 241
```

#Test error rate

```
svm_rad_test <- predict(radial_svm, oj_test)
table(svm_rad_test, oj_test$Purchase)
```

```
##
## svm_rad_test  CH  MM
##              CH 147  34
##              MM  17  72
```

The training error rate is 14.5% and the test error rate is 18.89%.

(g)

Polynomial Kernel

```
oj_poly <- svm(Purchase ~ ., data = oj_train, kernel = "polynomial", degree = 2)
summary(oj_poly)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "polynomial",
##      degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   1
##   degree:   2
##   gamma:   0.05555556
##   coef.0:   0
##
## Number of Support Vectors:  441
##
## ( 224 217 )
##
##
## Number of Classes:  2
```

```
##
## Levels:
## CH MM
```

The SVM with polynomial kernel creates 441 support vectors out of the 800 training observations. Out of the 441 support vectors, 224 belong to level CH and 217 to level MM.

```
#Training error rate
```

```
poly_train <- predict(oj_poly, oj_train)
table(poly_train, oj_train$Purchase)
```

```
##
## poly_train CH MM
## CH 452 107
## MM 37 204
```

```
#Test error rate
```

```
poly_test <- predict(oj_poly, oj_test)
table(poly_test, oj_test$Purchase)
```

```
##
## poly_test CH MM
## CH 149 46
## MM 15 60
```

```
(poly_error<- (107+37)/(452+107+37+204))
```

```
## [1] 0.18
```

```
(plove_error <- (46+15)/(149+46+15+60))
```

```
## [1] 0.2259259
```

The training error rate is 18% and the test error rate is 22.59%.

```
#For optimal cost
```

```
poly_tune <- tune(svm, Purchase ~ ., data = oj_train, kernel = "polynomial", degree = 2, ranges =
                  data.frame(cost = seq(0.01, 10, length.out = 25)))
summary(poly_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost
## 10
##
## - best performance: 0.1775
##
## - Detailed performance results:
## cost error dispersion
## 1 0.01000 0.38875 0.06192794
## 2 0.42625 0.20000 0.03996526
## 3 0.84250 0.19875 0.04143687
```

```
## 4    1.25875 0.19500 0.04338138
## 5    1.67500 0.19125 0.04411554
## 6    2.09125 0.19125 0.04372023
## 7    2.50750 0.19125 0.04604120
## 8    2.92375 0.18875 0.04839436
## 9    3.34000 0.19125 0.05138701
## 10   3.75625 0.19375 0.04868051
## 11   4.17250 0.19250 0.04866267
## 12   4.58875 0.19125 0.04641674
## 13   5.00500 0.19125 0.04896498
## 14   5.42125 0.19125 0.05138701
## 15   5.83750 0.18750 0.04526159
## 16   6.25375 0.18750 0.04526159
## 17   6.67000 0.18500 0.04556741
## 18   7.08625 0.18625 0.04767147
## 19   7.50250 0.18500 0.04851976
## 20   7.91875 0.18625 0.04693746
## 21   8.33500 0.18500 0.04816061
## 22   8.75125 0.18250 0.04901814
## 23   9.16750 0.18000 0.05006940
## 24   9.58375 0.18000 0.05006940
## 25  10.00000 0.17750 0.04816061
```

The optimal cost here is 10.

```
#Training error rate
poly_oj <- svm(Purchase ~ ., data = oj_train, kernel = "polynomial", cost = poly_tune$best.parameters$cost)
train_poly <- predict(poly_oj, oj_train)
table(train_poly, oj_train$Purchase)
```

```
##
## train_poly  CH  MM
##           CH 449  74
##           MM  40 237

(tr_err_poly <- (74+40)/(449+74+40+237))
```

```
## [1] 0.1425
```

```
#Test error rate
test_poly <- predict(poly_oj, oj_test)
table(test_poly, oj_test$Purchase)
```

```
##
## test_poly  CH  MM
##          CH 147  37
##          MM  17  69

(te_err_poly <- (37+17)/(147+37+17+69))
```

```
## [1] 0.2
```

The training error rate is 14.25% and the test error rate is 20%.

(h)

The support vector classifier or the SVM with linear kernel and with cost 3.75625 gives the best results in terms of the test error rate. It gives the smallest test error rate of 18.15%, among all the approaches.